

Practical work 06 – 21th of October 2025

Logistic Regression

Summary for the organisation :

- The practical works (PW) are optional. If you submit a PW, we provide **feedbacks** and you get potential **bonuses** on the final grade as explained in the first week. Submissions should be done before the date specified in Moodle.
- **Rule 1.** Submit 1 archive (*.zip) with your Python notebooks. Do not include datasets unless specific instructions, but do include all necessary files to reproduce your experiments.
- **Rule 2.** The archive file name must contain the number of the practical work, followed by the family names of the team members by alphabetical order, for example 02_dupont_muller_smith.zip. Put also the name of the team members in the body of the notebook (in first cell). Only one submission per team.
- **Rule 3.** We don't give bonuses for submissions that do not compile (missing files are a common source of errors...). So, make sure that your whole notebooks give the expected solutions by clearing all cells and running them all before submitting.

Exercise 1 Classification to predict student admission

This is a continuation of the exercise of previous week, where the objective is to build a classification system to predict whether a student gets admitted into a university or not based on their results on two exams¹. For each training example n , you have the applicants scores on two exams $(x_{n,1}, x_{n,2})$ and the admissions decision y_n . The dataset is illustrated on Figure 1.

a. Logistic regression classifier with linear decision boundary

The objective is to implement a classifier based on a logistic regression approach :

$$h_{\theta}(\mathbf{x}) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2)$$

1. Data source : Andrew Ng - Machine Learning class Stanford




FIGURE 1 – Training data

- a) In a similar way as for the exercise of the previous week, read the training data from file `student-dataset-train.csv`. The first two columns are x_1 and x_2 . The last column holds the class label y . Build the design matrix X as follow :

$$X = \begin{pmatrix} 1 & x_{1,1} & x_{1,2} \\ 1 & \vdots & \vdots \\ 1 & x_{N,1} & x_{N,2} \end{pmatrix} \quad (1)$$

Check that the shape of X is (100,3) and that the shape of y is (100,).

- b) Implement a z-norm normalization of the training set. You need to store the normalization values (μ, σ) for later as they will be needed to normalize the test set.
- c) Implement a sigmoid function $g(z) = \frac{1}{1+e^{-z}}$. Use numpy to compute the exp so that your function can take numpy arrays as input. Check your implementation by plotting the sigmoid function.
- d) Implement the hypothesis function $h_\theta(\mathbf{x})$. Hint : implement it so that the computation can take the full array X with $h(\mathbf{x})$ broadcasted to all training samples.
- e) Implement the objective function $J(\theta)$:

$$J(\theta) = \frac{1}{N} \sum_{n=1}^N y_n \log h_\theta(\mathbf{x}_n) + (1 - y_n) \log(1 - h_\theta(\mathbf{x}_n))$$

- f) In a similar way as in PW02 and PW03, implement the gradient ascent with the update rule :

$$\theta_i \leftarrow \theta_i + \alpha \frac{1}{N} \sum_{n=1}^N (y_n - h_\theta(\mathbf{x}_n)) x_{n,i}$$

- g) Test your implementation by running a gradient ascent. Hints : use a small α . e.g. 0.001, store the evolution of the objective function $J(\theta)$ during the epochs to make a plot, use a large number of epochs, e.g. 2000000.
- h) Compute the correct classification rate on `student-dataset-test.csv` after convergence as you have an estimator of the posterior probabilities with

$$P(y_n = 1 | \mathbf{x}_n; \theta) = h_\theta(\mathbf{x}_n)$$

$$P(y_n = 0 | \mathbf{x}_n; \theta) = 1 - h_\theta(\mathbf{x}_n)$$

This means that you can take the decisions $\hat{y}_n = 1$ if $h_\theta(\mathbf{x}_n) \geq 0.5$ and $\hat{y}_n = 0$ if $h_\theta(\mathbf{x}_n) < 0.5$.

- i) Draw the decision boundary of your system on top of the scatter plot of the testing data.
- j) Compare the performance of the logistic regression system with the ones of previous's week.

b. Optional - Stochastic gradient ascent

Redo the experiments of 2.a with a stochastic gradient ascent.

c. Logistic regression classifier with non-linear decision boundary

Redo the experiments of 2.a by increasing the complexity of the model in order to have a non-linear decision boundary :

$$h_\theta(\mathbf{x}) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1^2 + \theta_4 x_2^2 + \theta_5 x_1 x_2 + \dots)$$

d. Using SciKit Learn

Redo one of the exercise a. or c. using SciKit Learn.

- a) Read the documentation of the function `SGDClassifier()` available in the toolkit SciKit Learn. This function implements stochastic gradient descent training for different linear systems such as Logistic Regression. For a logistic regression, the `loss` parameter should be set to `'log'`.
- b) Use the `fit()` and `predict()` methods of this classifier on the student data.
- c) Compute the performances and compare it to your own implementations.

Exercise 2 Optional - Classification on MNIST

Using the SciKit Learn toolkit, train a multi-class logistic regression on the MNIST problem and compare the performances with the KNN of PW2.

For the “daring” ones, implement a multi-class version based on your own implementation and run it against MNIST data.

Exercise 3 Review questions

- a) Why do we have a gradient ascent in the case of logistic regression while we had a gradient descent with linear regression? Can we convert the gradient ascent of logistic regression into a gradient descent? If yes, how?
- b) Assuming a logistic regression with a linear decision boundary taking as input samples in two dimensions (x_1, x_2) , in which case do we get 0.5 as output of the classification system? Express your answer with an equation.
- c) What is the computational trick to avoid numerical problems in the computation of $J(\theta)$ for the logistic regression? In which situations (for what type of inputs) do we risk to observe such numerical problems?
- d) A logistic regression can classify between 2 classes. How can we build a multi-class (with K classes) system with logistic regression?